



EIGRP IPv6 & Routing

Jirasak Sittigorn

Internetworking Standards & Technologies

Department of Computer Engineering, Faculty of Engineering
King Mongkut's Institute of Technology Ladkrabang

Cisco | Networking Academy®
Mind Wide Open™

EIGRP

- EIGRP Characteristics

- Explain the features and characteristics of EIGRP.
- Describe the basic features of EIGRP.
- Describe the types of packets used to establish and maintain an EIGRP neighbor adjacency.
- Describe the encapsulation of an EIGRP messages.

- Implement EIGRP for IPv4

- Implement EIGRP for IPv4 in a small to medium-sized business network.
- Configure EIGRP for IPv4 in a small routed network.
- Verify EIGRP for IPv4 operation in a small routed network.

- EIGRP Operation

- Explain how EIGRP operates in a small to medium-sized business network.
- Explain how EIGRP forms neighbor relationships.
- Explain the metrics used by EIGRP.
- Explain how DUAL operates and uses the topology table.
- Describe events that trigger EIGRP updates.

		Interior Gateway Protocols			Exterior Gateway Protocols	
		Distance Vector Routing Protocols		Link State Routing Protocols		Path Vector
Classful	RIP	IGRP			EGP	
	RIPv2	EIGRP	OSPFv2		IS-IS	BGPv4
	RIPng	EIGRP for IPv6	OSPFv3		IS-IS for IPv6	BGPv4 for IPv6

EIGRP Characteristics

- EIGRP Basic Features

- Enhanced IGRP is a Cisco-proprietary *distance-vector routing protocol* released in 1992.

- EIGRP was created as a classless version of IGRP.
- Ideal choice for large, multiprotocol networks built primarily on Cisco routers.



EIGRP Feature	Description
Diffusing Update Algorithm (DUAL)	<ul style="list-style-type: none"> • EIGRP uses DUAL as its routing algorithm. • DUAL guarantees loop-free and backup paths throughout the routing domain.
Establishing Neighbor Adjacencies	<ul style="list-style-type: none"> • EIGRP establishes relationships with directly connected EIGRP routers. • Adjacencies are used to track the status of these neighbors.
Reliable Transport Protocol	<ul style="list-style-type: none"> • EIGRP RTP provides delivery of EIGRP packets to neighbors. • RTP and neighbor adjacencies are used by DUAL.
Partial and Bounded updates	<ul style="list-style-type: none"> • Instead of periodic updates, EIGRP sends partial triggered updates when a path or metric changes. • Only those routers that require the information are updated minimizing bandwidth use.
Equal and Unequal Cost Load Balancing	<ul style="list-style-type: none"> • EIGRP supports equal cost load balancing and unequal cost load balancing, which allows administrators to better distribute traffic flow in their networks.

EIGRP Characteristics (Basic Features)


- EIGRP uses protocol-dependent modules (PDMs) to support different protocols such as IPv4, IPv6, and legacy protocols IPX and AppleTalk.
- PDMs are responsible for:
 - Maintaining EIGRP neighbor and topology tables
 - Computing the metric using DUAL
 - Interfacing DUAL and routing table
 - Implementing filtering and access lists
 - Performing redistribution with other routing protocols

EIGRP maintains individual tables for each routed protocol.




EIGRP Characteristics (Basic Features)

- RTP is the EIGRP Transport layer protocol used for the delivery and reception of EIGRP packets.
- Not all RTP packets are sent reliably.
 - Reliable packets require explicit acknowledgement from destination
 - Update, Query, Reply
 - Unreliable packets do not require acknowledgement from destination
 - Hello, ACK



EIGRP Characteristics (Basic Features)

- EIGRP supports authentication and is recommended.
 - EIGRP authentication ensures that routers only accept routing information from other routers that have been configured with the same password or authentication information.
- Note:
 - Authentication does not encrypt the EIGRP routing updates.



EIGRP Characteristics


- EIGRP Packet Types
 - IP EIGRP relies on 5 types of packets to maintain its various tables and establish complex relationships with neighbor routers.

EIGRP Packet Types

Packet Type	Used to...
Hello	Discover other EIGRP routers in the network.
Update	Convey routing information to known destinations.
Acknowledgement	Acknowledge the receipt of any EIGRP packet.
Query	Request specific information from a neighbor router.
Reply	Respond to a query.

EIGRP Characteristics (Packet Types)

- Hello packets are used to discover & form adjacencies with neighbors.
 - On hearing Hellos, a router creates a neighbor table and the continued receipt of Hellos maintains the table.
- Hello packets are always sent unreliably.
 - Therefore Hello packets do not require acknowledgment.




Bandwidth	Example Link	Default Hello Interval	Default Hold Time
1.544 Mb/s	Multipoint Frame Relay	60 seconds	180 seconds
Greater than 1.544 Mb/s	T1, Ethernet	5 seconds	15 seconds

EIGRP uses multicast and unicast rather than broadcast.

- As a result, end stations are unaffected by routing updates or queries.
- The EIGRP multicast IPv4 address is **224.0.0.10**
- The EIGRP multicast IPv6 address is **FF02::A**.


EIGRP Characteristics (Packet Types)

- EIGRP Update packets are used to propagate routing information.
 - Sent to initially exchange topology information or topology change.
 - EIGRP updates only contain needed routing information and are unicast to routers that require it.
 - Update packets are sent reliably and therefore requires acknowledgements.
- Acknowledgements packets are “dataless” Hello packets used to indicate receipt of any EIGRP packet during a “reliable” (i.e., RTP) exchange.
 - Used to acknowledge the receipt of Update packets, Query packets, and Reply packets.




EIGRP Characteristics (Packet Types)

- Query and reply packets are used by DUAL when searching for networks.
- They both use reliable delivery and therefore require acknowledgement.
- Queries can use multicast or unicast, whereas Replies are always sent as unicast.




EIGRP Characteristics (Messages)

- EIGRP Messages
 - EIGRP frame contains destination multicast address 01-00-5E-00-00-0A.
 - The IP packet header contains destination IP address 224.0.0.10 and identifies this packet as an EIGRP packet (protocol 88).
 - The data portion of the EIGRP message includes:
 - **Packet header** - The EIGRP packet header identifies the type of EIGRP message.
 - **Type/Length/Value (TLV)** - The TLV field contains EIGRP parameters, IP internal and external routes.
 - EIGRP for IPv6 is encapsulated using an IPv6 header with multicast address FF02::A and the next header field set to protocol 88.




EIGRP Characteristics (Messages)


- EIGRP messages include the header with an Opcode field that specifies the type of EIGRP packet (Hello, Ack, Update, Query, and Reply) and the AS number field.




EIGRP TLV: EIGRP Parameters



EIGRP TLV: Internal Routes




EIGRP TLV: External Routes




Implement EIGRP for IPv4

- Configure EIGRP with IPv4
 - The routers in the topology have a starting configuration that includes addresses on the interfaces. There is currently no static routing or dynamic routing configured on any of the routers.



Implement EIGRP for IPv4 (Configure)

- An Autonomous System (AS) is a collection of networks under the control of a single authority (reference RFC 1930).
 - AS numbers are needed to exchange routes between AS.
 - AS numbers are managed by IANA and assigned by RIRs to ISPs, Internet Backbone providers, and institutions connecting to other institutions using AS numbers.
- AS numbers are usually 16-bit numbers, ranging from 0 to 65535.
 - Since 2007, AS numbers can now be 32 bits, therefore increasing the number of AS numbers to over 4 billion.




Implement EIGRP for IPv4 (Configure)

- To configure EIGRP, use the **router eigrp AS-#** command.
 - The **AS-#** functions as a process ID.
 - The AS number used for EIGRP configuration is only significant to the EIGRP routing domain.
 - All routers in the EIGRP routing domain must use the same AS number (process ID number).
- Note:
 - Do NOT configure multiple instances of EIGRP on the same router.

Implement EIGRP for IPv4 (Configure)

- The EIGRP router ID is used to uniquely identify each router in the EIGRP routing domain.
- Routers use the following three criteria to determine its router ID:
 - Use the address configured with the **eigrp router-id** **ipv4-address** router config command.
 - If the router ID is not configured, choose the highest IPv4 address of any of its loopback interfaces.
 - If no loopback interfaces are configured, choose the highest active IPv4 address of any of its physical interfaces.



Implement EIGRP for IPv4 (Configure)

```
R1(config)# router eigrp 1
R1(config-router)# eigrp router-id 1.1.1.1
R1(config-router)#

```

```
R2(config)# router eigrp 1
R2(config-router)# eigrp router-id 2.2.2.2
R2(config-router)#


```

```
R1# show ip protocols
*** IP Routing is NSF aware ***

Routing Protocol is "eigrp 1"
  Outgoing update filter list for all interfaces is not
    set
  Incoming update filter list for all interfaces is not
    set
  Default networks flagged in outgoing updates
  Default networks accepted from incoming updates
  EIGRP-IPv4 Protocol for AS(1)
    Metric weight K1=1, K2=0, K3=1, K4=0, K5=0
    NSF-aware route hold timer is 240
    Router-ID: 1.1.1.1
    Topology : 0 (base)
      Active Timer: 3 min
      Distance: internal 90 external 170
      Maximum path: 4
      Maximum hopcount 100
      Maximum metric variance 1

  Automatic Summarization: disabled
  Maximum path: 4
  Routing for Networks:
  Routing Information Sources:
    Gateway          Distance      Last Update
    Distance: internal 90 external 170

```



Implement EIGRP for IPv4 (Configure)

- Use the **network network-number [wildcard-mask]** router config command to enable and advertise a network in EIGRP.
 - It enables the interfaces configured for that network address to begin transmitting & receiving EIGRP updates
 - Includes network or subnet in EIGRP updates

```
Enables EIGRP for the interfaces on subnets in 172.16.1.0/24 and 172.16.3.0/30.

R1(config)# router eigrp 1
R1(config-router)# network 172.16.0.0
R1(config-router)# network 192.168.10.0
R1(config-router)#

Enables EIGRP for the interfaces on subnet 192.168.10.4/30.
```

Implement EIGRP for IPv4 (Configure)

- A wildcard mask is similar to a subnet mask but is calculated by subtracting a SNM from 255.255.255.255.
- For example, if the SNM is 255.255.255.252:
 - 255.255.255.255
 - 255.255.255.252
 - 0. 0. 0. 3 Wildcard mask
- EIGRP also automatically converts a subnet mask to its wildcard mask equivalent.
 - E.g., entering 192.168.10.8 255.255.255.252 automatically converts to 192.168.10.8 0.0.0.3

```
R2(config)# router eigrp 1
R2(config-router)# network 192.168.10.8 0.0.0.3
R2(config-router)
```

```
R2(config)# router eigrp 1
R2(config-router)# network 192.168.10.8 255.255.255.252
R2(config-router)# end
R2# show running-config | section eigrp 1
router eigrp 1
network 172.16.0.0
network 192.168.10.8 0.0.0.3
eigrp router-id 2.2.2.2
R2#
```

Implement EIGRP for IPv4 (Configure)

- Passive interfaces prevent EIGRP updates out a specified router interface.

Router (config-router) #passive-interface type number [default]

- Set a particular interface or all router interfaces to passive.
 - The **default** option sets all router interfaces to passive.
 - Prevents neighbor relationships from being established.
 - Routing updates from a neighbor are ignored.

```
R3(config)# router eigrp 1
R3(config-router)# passive-interface gigabitethernet 0/0
```

```
R3# show ip protocols
*** IP Routing is NSF aware ***

Routing Protocol is "eigrp 1"
<output omitted>
Routing for Networks:
  192.168.1.0
  192.168.10.4/30
  192.168.10.8/30
Passive Interface(s):
  GigabitEthernet0/0
Routing Information Sources:
  Gateway          Distance      Last Update
  192.168.10.5        90          01:37:57
  192.168.10.9        90          01:37:57
  Distance: internal 90 external 170
R3#
```

Implement EIGRP for IPv4

- Verify EIGRP with IPv4
 - Use the `show ip eigrp neighbors` command to view the neighbor table and verify that EIGRP has established an adjacency with its neighbors.
 - The output displays a list of each adjacent neighbor.

```
R1# show ip eigrp neighbors
EIGRP-IPv4 Neighbors for AS(1)
  H   Address           Interface      Hold  Uptime       SRTT    RTO     Q      Seq
  1  192.168.10.6      Se0/0/1        11   04:57:14    27    162     0      8
  0  172.16.3.2        Se0/0/0        13   07:53:46    20    120     0      10
R1#
```

The diagram shows the output of the `show ip eigrp neighbors` command on router R1. The output lists two neighbors: one with address 192.168.10.6 and another with address 172.16.3.2. Arrows point from the callout boxes to the corresponding columns in the command output.

H	Address	Interface	Hold (sec)	Uptime	SRTT (ms)	RTO	Q	Seq Num
1	192.168.10.6	Se0/0/1	11	04:57:14	27	162	0	8
0	172.16.3.2	Se0/0/0	13	07:53:46	20	120	0	10

Callout boxes and their descriptions:

- Neighbor's IPv4 Address: Points to the "Address" column.
- Local Interface receiving EIGRP Hello packets: Points to the "Interface" column.
- Seconds remaining before declaring neighbor down: Points to the "Hold (sec)" column.
- Amount of time since this neighbor was added to the neighbor table: Points to the "Uptime" column.

A note at the bottom of the callout for "Seconds remaining before declaring neighbor down" states: "The current hold time is reset to the maximum hold time whenever a Hello packet is received".

Implement EIGRP for IPv4 (Verify)

- The **show ip protocols** command is useful to identify the parameters and other information about the current state of any active IPv4 routing protocol processes configured on the router.
- For example, in the command output in the figure:
- EIGRP is an active dynamic routing protocol on R1 configured with the autonomous system number 1.
 1. The EIGRP router ID of R1 is 1.1.1.1.
 2. The EIGRP administrative distances on R1 are internal AD of 90 and external of 170 (default values).
 3. By default, EIGRP does not automatically summarize networks. Subnets are included in the routing updates.
 4. The EIGRP neighbor adjacencies R1 has with other routers used to receive EIGRP routing updates.

```
R1# show ip protocols
*** IP Routing is NSF aware ***

Routing Protocol is "eigrp 1" ① Routing protocol and Process ID (AS Number)
Outgoing update filter list for all interfaces is not set
Incoming update filter list for all interfaces is not set
Default networks flagged in outgoing updates
Default networks accepted from incoming updates
EIGRP-IPv4 Protocol for AS(1)
  Metric weight K1=1, K2=0, K3=1, K4=0, K5=0
  NSF-aware route hold timer is 240
  Router-ID: 1.1.1.1 ② EIGRP Router ID

  Topology : 0 (base)
  Active Timer: 3 min
  Distance: internal 90 external 170 ③ EIGRP Administrative Distances

  Maximum path: 4
  Maximum hopcount 100
  Maximum metric variance 1

  Automatic Summarization: disabled ④ EIGRP Automatic Summarization is disabled.


  Maximum path: 4
  Routing for Networks:
    172.16.0.0
    192.168.10.0
  Routing Information Sources: ⑤ EIGRP Routing Information Sources lists all the EIGRP routing sources the IOS uses to build its IPv4 routing table.

    Gateway      Distance   Last Update
    192.168.10.6        90   00:40:20
    172.16.3.2          90   00:40:20

  Distance: internal 90 external 170
```

R1#

Implement EIGRP for IPv4 (Verify)



```
R1# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
<output omitted>
```

Gateway of last resort is not set

```

172.16.0.0/16 is variably subnetted, 5 subnets, 3 masks
C 172.16.1.0/24 is directly connected, GigabitEthernet0/0
L 172.16.1.1/32 is directly connected, GigabitEthernet0/0
D 172.16.2.0/24 [90/2170112] via 172.16.3.2, 00:14:35, Serial0/0/0
C 172.16.3.0/30 is directly connected, Serial0/0/0
L 172.16.3.1/32 is directly connected, Serial0/0/0
D 192.168.1.0/24 [90/2170112] via 192.168.10.6, 00:13:57, Serial0/0/1
  192.168.10.0/24 is variably subnetted, 3 subnets, 2 masks
C 192.168.10.4/30 is directly connected, Serial0/0/1
L 192.168.10.5/32 is directly connected, Serial0/0/1
D 192.168.10.8/30 [90/2681856] via 192.168.10.6, 00:50:42, Serial0/0/1
  [90/2681856] via 172.16.3.2, 00:50:42, Serial0/0/0
  
```

R1#

```
R2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
<output omitted>
Gateway of last resort is not set

172.16.0.0/16 is variably subnetted, 5 subnets, 3 masks
D 172.16.1.0/24 [90/2170112] via 172.16.3.1, 00:11:05, Serial0/0/0
C 172.16.2.0/24 is directly connected, GigabitEthernet0/0
L 172.16.2.1/32 is directly connected, GigabitEthernet0/0
C 172.16.3.0/30 is directly connected, Serial0/0/0
L 172.16.3.2/32 is directly connected, Serial0/0/0
D 192.168.1.0/24 [90/2170112] via 192.168.10.10, 00:15:16, Serial0/0/1
  192.168.10.0/24 is variably subnetted, 3 subnets, 2 masks
D 192.168.10.4/30 [90/2681856] via 192.168.10.10, 00:52:00, Serial0/0/1
  [90/2681856] via 172.16.3.1, 00:52:00, Serial0/0/0
C 192.168.10.8/30 is directly connected, Serial0/0/1
L 192.168.10.9/32 is directly connected, Serial0/0/1
  209.165.200.0/24 is variably subnetted, 2 subnets, 2 masks
C 209.165.200.224/27 is directly connected, Loopback209
L 209.165.200.225/32 is directly connected, Loopback209
R2#
```

```
R3# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
<output omitted>
```

Gateway of last resort is not set

```


172.16.0.0/16 is variably subnetted, 3 subnets, 2 masks
D 172.16.1.0/24 [90/2170112] via 192.168.10.5, 00:12:00, Serial0/0/0
D 172.16.2.0/24 [90/2170112] via 192.168.10.9, 00:16:49, Serial0/0/1
D 172.16.3.0/30 [90/2681856] via 192.168.10.9, 00:52:55, Serial0/0/1
  [90/2681856] via 192.168.10.5, 00:52:55, Serial0/0/0
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C 192.168.1.0/24 is directly connected, GigabitEthernet0/0
L 192.168.1.1/32 is directly connected, GigabitEthernet0/0
  192.168.10.0/24 is variably subnetted, 4 subnets, 2 masks
C 192.168.10.4/30 is directly connected, Serial0/0/0
L 192.168.10.6/32 is directly connected, Serial0/0/0
C 192.168.10.8/30 is directly connected, Serial0/0/1
L 192.168.10.10/32 is directly connected, Serial0/0/1
  
```

R3#

EIGRP Operation

- EIGRP Initial Route Discovery


1. Router R1 starts has joined the EIGRP routing domain and sends an EIGRP Hello packet out all EIGRP enabled interfaces.
2. Router R2 receives the Hello packet and adds R1 to its neighbor table.
 - R2 sends an Update packet that contains all the routes it knows.
 - R2 also sends an EIGRP Hello packet to R1.



3. R1 updates its neighbor table with R2.
 - After both routers have exchanged Hellos, the neighbor adjacency is established.

EIGRP Operation (Initial Route Discovery)


1. R1 adds all update entries from R1 to its topology table.
 - The topology table includes all destinations advertised by neighboring (adjacent) routers and the cost (metric) to reach each network.
2. EIGRP update packets use reliable delivery; therefore, R1 replies with an EIGRP acknowledgment packet informing R2 that it has received the update.
3. R1 sends an EIGRP update to R2 advertising the routes that it is aware of, except those learned from R2 (split horizon).
4. R2 receives the EIGRP update from R1 and adds this information to its own topology table.
5. R2 responds to R1's EIGRP update packet with an EIGRP acknowledgment.



EIGRP Operation (Initial Route Discovery)

1. R1 uses DUAL to calculate the best routes to each destination, including the metric and the next-hop router and updates its routing table with the best routes.

2. Similarly, R2 uses DUAL and updates its routing table with the best newly discovered routes.



- At this point, EIGRP on both routers is considered to be in the converged state.

EIGRP Operation

- EIGRP Metrics
 - EIGRP uses a composite metric which can be based on the following metrics:
 - Bandwidth: The lowest bandwidth between source and destination.
 - Delay: The cumulative interface delay along the path
 - Reliability: (Optional) Worst reliability between source and destination.
 - Load: (Optional) Worst load on a link between source and destination.
 - The EIGRP composite metric formula consists metric weights with values K1 to K5.
 - K1 represents bandwidth, K3 delay, K4 load, and K5 reliability.

Default Values:
 K1 (bandwidth) = 1
 K2 (load) = 0
 K3 (delay) = 1
 K4 (reliability) = 0
 K5 (reliability) = 0

Note.

- It is often incorrectly stated that EIGRP can also use the smallest MTU in the path.

Default Composite Formula:

`metric = [K1*bandwidth + K3*delay] * 256`

Complete Composite Formula:

`metric = [K1*bandwidth + (K2*bandwidth) / (256 - load) + K3*delay] * [K5 / (reliability + K4)]`

(Not used if "K" values are 0)

Note: This is a conditional formula. If K5 = 0, the last term is replaced by 1 and the formula becomes: Metric = [K1*bandwidth + (K2*bandwidth) / (256-load) + K3*delay] * 256

```
Router(config-router)# metric weights tos k1 k2 k3 k4 k5
```

EIGRP Operation (Metrics)

– Use the **show interfaces** command to examine the values used for bandwidth, delay, reliability, and load.

- **BW** - Bandwidth of the interface (in kb/s).
- **DLY** - Delay of the interface (in microseconds).
- **Reliability** - Reliability of the interface as a fraction of 255 (255/255 is 100% reliability).
- **Txload, Rxload** - Transmit and receive load on the interface as a fraction of 255 (255/255 is completely saturated), calculated as an exponential average over five minutes.

```
R1# show interfaces serial 0/0/0
Serial0/0/0 is up, line protocol is up
  Hardware is WIC MBRD Serial
  Internet address is 172.16.3.1/30
  MTU 1500 bytes, BW 1544 Kbit/sec, DLY 20000 usec,
    reliability 255/255, txload 1/255, rxload 1/255
  Encapsulation HDLC, loopback not set
<output omitted>
R1#  
  
R1# show interfaces gigabitethernet 0/0
GigabitEthernet0/0 is up, line protocol is up
  Hardware is CN Gigabit Ethernet, address is fc99.4775.c3e0
  (bia fc99.4775.c3e0)
  Internet address is 172.16.1.1/24
  MTU 1500 bytes, BW 100000 Kbit/sec, DLY 100 usec,
    reliability 255/255, txload 1/255, rxload 1/255
  Encapsulation ARPA, loopback not set
<output omitted>
R1#
```

EIGRP Operation (Metrics)

- Use the following interface configuration mode command to modify the bandwidth metric:


```
Router (config-if) #
bandwidth kilobits-bandwidth-value
```

- Use the **show interfaces** command to verify the new bandwidth parameters.

```
R1# show interface s 0/0/0
Serial0/0/0 is up, line protocol is up
  Hardware is WIC MBRD Serial
  Internet address is 172.16.3.1/30
  MTU 1500 bytes, BW 64 Kbit/sec, DLY 20000 usec,
    reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R1#
```

```
R2# show interface s 0/0/0
Serial0/0/0 is up, line protocol is up
  Hardware is WIC MBRD Serial
  Internet address is 172.16.3.2/30
  MTU 1500 bytes, BW 64 Kbit/sec, DLY 20000 usec,
    reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R2#
```

```
R2(config)# interface s 0/0/0
R2(config-if)# bandwidth 64
R2(config-if)# exit
R2(config)# interface s 0/0/1
R2(config-if)# bandwidth 1024
```



```
R1(config)# interface s 0/0/0
R1(config-if)# bandwidth 64
```

```
R3(config)# interface s 0/0/1
R3(config-if)# bandwidth 1024
```

EIGRP Operation (Metrics)

- Delay is a measure of the time it takes for a packet to traverse a route.
- The delay (DLY) metric is not measured dynamically.
 - It is a static value measured in microseconds (μ s or usec) based on the type of link to which the interface is connected.
- The delay value is calculated using the cumulative (sum) of all interface delays along the path, divided by 10.

Media	Delay In usec
Gigabit Ethernet	10
Fast Ethernet	100
FDDI	100
16M Token Ring	630
Ethernet	1,000
T1 (Serial Default)	20,000
DS0 (64 Kbps)	20,000
1024 Kbps	20,000
56 Kbps	20,000

EIGRP Operation (Metrics)

– We can determine the EIGRP metric as follows:

1. Determine the link with the slowest bandwidth and use that value to calculate bandwidth ($10,000,000/\text{bandwidth}$).
2. Determine the delay value for each outgoing interface on the way to the destination and add the delay values and divide by 10 (sum of delay/10).
3. This composite metric produces a 24-bit value which EIGRP multiplies with 256.

$$[K1 * \text{bandwidth} + K3 * \text{delay}] * 256 = \text{Metric}$$

Because K1 and K3 both equal 1, the formula becomes:

$$(\text{Bandwidth} + \text{Delay}) * 256 = \text{Metric}$$

$$((10,000,000 / \text{bandwidth}) + (\text{sum of delay} / 10)) * 256 = \text{Metric}$$

```
R2# show ip route
```

```
D 192.168.1.0/24 [90/3012096] via 192.168.10.10, 00:12:32, Serial0/0/1
```

EIGRP Operation (Metrics)

- How does EIGRP determine the following metric?

R2# show ip route

<output omitted>

D 192.168.1.0/24 [90/3012096] via 192.168.10.10, 00:12:32, Serial0/0/1

- EIGRP Composite Metric = (Bandwidth + Delay) × 256

▪ Bandwidth = 10,000,000 / slowest bandwidth

```
R2# show interface s 0/0/1
Serial0/0/1 is up, line protocol is up
Hardware is WIC MBRD Serial
Internet address is 192.168.10.9/30
MTU 1500 bytes, BW 1024 Kbit/sec, DLY 20000 usec,
reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R2#
```

```
R3# show interface g 0/0
GigabitEthernet0/0 is up, line protocol is up
Hardware is CN Gigabit Ethernet, address is fc99.4771.7a20 (bia fc99.4771.7a20)
Internet address is 192.168.1.1/24
MTU 1500 bytes, BW 1000000 Kbit/sec, DLY 10 usec,
reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R3#
```

▪ Bandwidth = 10,000,000 / 1024 = 9765

▪ Delay = (Sum of all delays) / 10

```
R2# show interface s 0/0/1
Serial0/0/1 is up, line protocol is up
Hardware is WIC MBRD Serial
Internet address is 192.168.10.9/30
MTU 1500 bytes, BW 1024 Kbit/sec, DLY 20000 usec,
reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R2#
```

```
R3# show interface g 0/0
GigabitEthernet0/0 is up, line protocol is up
Hardware is CN Gigabit Ethernet, address is fc99.4771.7a20 (bia fc99.4771.7a20)
Internet address is 192.168.1.1/24
MTU 1500 bytes, BW 1000000 Kbit/sec, DLY 10 usec,
reliability 255/255, txload 1/255, rxload 1/255
<output omitted>
R3#
```

▪ Delay = (20,000 + 10) / 1024 = 2001

- EIGRP Composite Metric = (9765 + 2001) × 256 = 3,012,096

EIGRP Operation

- DUAL and the Topology Table
 - EIGRP uses the Diffusing Update Algorithm (DUAL) to provide the best and backup loop-free paths.
 - DUAL uses several terms, which are discussed in more detail throughout this section:


Term	Description
Successor	<ul style="list-style-type: none">• Is a neighboring router that is used for packet forwarding and is the least-cost route to the destination network.• The IP address of a successor is shown in a routing table entry right after the word “via”.
Feasible Successors (FS)	<ul style="list-style-type: none">• These are the “Backup paths” that are a loop-free.• Must comply to a feasibility condition.
Reported Distance (RD)	<ul style="list-style-type: none">• Also called “advertised distance”, this is the reported metric from the neighbor advertising the route.• If the RD metric is less than the FD, then the next-hop router is downstream and there is no loop.
Feasible Distance (FD)	<ul style="list-style-type: none">• This is the actual metric of a route from the current router.• Is the lowest calculated metric to reach the destination network.• FD is the metric listed in the routing table entry as the second number inside the brackets.

EIGRP Operation (DUAL-Topology)

- Routing loops, even temporary ones, can be detrimental to network performance and EIGRP prevents routing loops with the DUAL algorithm.
 - The DUAL algorithm is used to obtain loop-freedom at every instance throughout a route computation.
- The decision process for all route computations is done by the DUAL Finite State Machine (FSM). An FSM is a workflow model, similar to a flow chart, which is composed of the following:
 - A finite number of stages (states)
 - Transitions between those stages
 - Operations
- The DUAL FSM tracks all routes and uses EIGRP metrics to select efficient, loop-free paths, and to identify the routes with the least-cost path to be inserted into the routing table.

EIGRP Operation (DUAL-Topology)

- A successor is a neighboring router with the least-cost route to the destination network.
 - The successor IP address is shown right after “via”.
 - FD is the lowest calculated metric to reach the destination network.
 - FD is the second number inside the brackets.
 - Also known as the “metric” for the route.
 - Notice that EIGRP’s best path for the 192.168.1.0/24 network is through router R3, and that the feasible distance is 3,012,096.






```
R2# show ip route  
<output omitted>  
D 192.168.1.0/24 [90/3012096] via 192.168.10.10, 00:12:32,  
Serial0/0/1
```

Feasible Distance

Successor


EIGRP Operation (DUAL-Topology)

- DUAL converges quickly because it can use backup paths known as Feasible Successors (FSs).
- A FS is a neighbor with a loop-free backup path to the same network as the successor.
 - A FS must satisfy the Feasibility Condition (FC).
 - The FC is met when a neighbor's Reported Distance (RD) is less than the local router's feasible distance.
 - If the reported distance is less, it represents a loop-free path.
- E.g., the RD of R1 (2,170,112) is less than R2's own FD (3,012,096) and therefore, R1 meets the FC and becomes the FS for R2 to the 192.168.1.0/24 network.



EIGRP Operation (DUAL-Topology)

- Topology table stores the following information required by DUAL to calculate distances and vectors to destinations.
 - The **reported distance (RD)** that each neighbor advertises for each destination
 - The **feasible distance (FD)** that this router would use to reach the destination via that neighbor.
- Use the **show ip eigrp topology** command to list all successors and FSs to destination networks.
 - Only the successor is installed into the IP routing table.
 - **Passive State** - Route is in stable state and available for use.
 - **Active State** - Route is being recomputed by DUAL.



```
R2# show ip eigrp topology
EIGRP-IPv4 Topology Table for AS(1)/ID(2.2.2.2)
Codes: P - Passive, A - Active, U - Update, Q - Query, R - Reply,
       r - reply Status, s - sia Status


P 172.16.2.0/24, 1 successors, FD is 2816
      via Connected, GigabitEthernet0/0
P 192.168.10.4/30, 1 successors, FD is 3523840
      via 192.168.10.10 (3523840/2169856), Serial0/0/1
      via 172.16.3.1 (41024000/2169856), Serial0/0/0
P 192.168.1.0/24, 1 successors, FD is 3012896
      via 192.168.10.10 (3012096/2816), Serial0/0/1
      via 172.16.3.1 (41024256/2170112), Serial0/0/0
P 172.16.3.0/30, 1 successors, FD is 40512000
      via Connected, Serial0/0/0
P 172.16.1.0/24, 1 successors, FD is 3524096
      via 192.168.10.10 (3524096/2170112), Serial0/0/1
      via 172.16.3.1 (40512256/2816), Serial0/0/0
P 192.168.10.8/30, 1 successors, FD is 3011840
      via Connected, Serial0/0/1

R2#
```

EIGRP Operation (DUAL-Topology)

- The first line in the topology table displays:

- P - Route in the passive state (the route is in a stable mode). If DUAL recalculates or searches for a new path, the route is in an active state and displays an A.
- 192.168.1.0/24** - Destination network is also found in the routing table.
- 1 successors - Displays the number of successors for this network. If there are multiple equal cost paths to this network, there are multiple successors.
- FD is 3012096** - FD, the EIGRP metric to reach the destination network. This is the metric displayed in the IP routing table.



```
R2# show ip eigrp topology
EIGRP-IPv4 Topology Table for AS(1)/ID(2.2.2.2)
Codes: P - Passive, A - Active, U - Update, Q - Query, R - Reply,
       r - reply Status, s - sia Status

P 172.16.2.0/24, 1 successors, FD is 2816
      via Connected, GigabitEthernet0/0
P 192.168.10.4/30, 1 successors, FD is 3523840
      via 192.168.10.10 (3523840/2169856), Serial0/0/1
      via 172.16.3.1 (41024000/2169856), Serial0/0/0
P 192.168.1.0/24, 1 successors, FD is 3012096
      via 192.168.10.10 (3012096/2816), Serial0/0/1
      via 172.16.3.1 (41024256/2170112), Serial0/0/0
P 172.16.3.0/30, 1 successors, FD is 40512000
      via Connected, Serial0/0/0
P 172.16.1.0/24, 1 successors, FD is 3524096
      via 192.168.10.10 (3524096/2170112), Serial0/0/1
      via 172.16.3.1 (40512256/2816), Serial0/0/0
P 192.168.10.8/30, 1 successors, FD is 3011840
      via Connected, Serial0/0/1
```

R2#

EIGRP Operation (DUAL-Topology)

- The partial output of the **show ip route** command displays the 192.168.1.0/24 route with the successor is R3 via 192.168.10.6 with an FD of 2,170,112.
- The **show ip eigrp topology** command only shows the successor 192.168.10.6, which is R3.
 - Notice there are no FSs.
- The **show ip eigrp topology all-links** command shows all possible paths to a network, including successors, FSs, and even those routes that are not FSs.

```
R1# show ip route
<output omitted>
D 192.168.1.0/24 [90/2170112] via 192.168.10.6,
01:23:13, Serial0/0/1
```

Feasible Distance
Next-hop router (R3) is the successor.

```
R1# show ip eigrp topology
<output omitted>
P 192.168.1.0/24, 1 successors, FD is 2170112
via 192.168.10.6 (2170112/2816), Serial0/0/1
```


Successor

```
R1# show ip eigrp topology all-links
<output omitted>
P 192.168.1.0/24, 1 successors, FD is 2170112, serno 9
via 192.168.10.6 (2170112/2816), Serial0/0/1
via 172.16.3.2 (41024256/3012096), Serial0/0/0
```

R1's Feasible Distance
R2's Reported Distance
Successor
Not a feasible successor


EIGRP Operation

- DUAL and Convergence
 - The DUAL Finite State Machine (FSM) contains all of the logic used to calculate and compare routes in an EIGRP network.
 - An FSM is an abstract machine, that defines a set of possible states that something can go through, what events cause those states, and what events result from those states.
 - Designers use FSMs to describe how a device, computer program, or routing algorithm reacts to a set of input events.



EIGRP Operation (DUAL-Convergence)

- IP EIGRP



EIGRP Operation (DUAL-Convergence)


- If the path to the successor fails and there are no FSs, DUAL puts the network into the active state and actively queries its neighbors for a new successor.
 - DUAL sends EIGRP queries asking other routers for a path to the network.
 - Other routers return EIGRP replies, letting the sender of the EIGRP query know that they have a path to the requested network. If there is no reply, the sender of the query does not have a route to this network.
 - If the sender receives EIGRP replies with a path to the requested network, the preferred path is added as the new successor and also added to the routing table.

IPV6 NETWORK ADDRESSES

- IPv4 Issues
 - The Need for IPv6
 - IPv4 and IPv6 Coexistence
- IPv6 Addressing
 - IPv6 Address Representation
 - Rule 1 - Omit Leading 0s
 - Rule 2 - Omit All 0 Segments
- Types of IPv6 Addresses
 - IPv6 Address Types
 - IPv6 Prefix Length
 - IPv6 Unicast Addresses
 - IPv6 Link-Local Unicast Addresses
- IPv6 Unicast Addresses
 - Structure of an IPv6 Global Unicast Address
 - Static Configuration of a Global Unicast Address
 - Dynamic Configuration - SLAAC *
 - Dynamic Configuration - DHCPv6 *
 - EUI-64 Process and Randomly Generated *
 - Dynamic Link-Local Addresses *
 - Static Link-Local Addresses *
 - Verifying IPv6 Address Configuration *




IPv4 Issues

- The Need for IPv6
 - IPv6 versus IPv4:
 - Has a larger 128-bit address space
 - 340 undecillion addresses
 - Solves limitations with IPv4
 - Adds enhancement like address auto-configuration.
 - Why IPv6 is needed:
 - Rapidly increasing Internet population
 - Depletion of IPv4
 - Issues with NAT
 - Internet of Things



IPv4 Issues

- IPv4 and IPv6 Coexistence
 - Migration from IPv4 to IPv6 Techniques
 - Dual stack - Devices run both IPv4 and IPv6 protocol stacks simultaneously.
 - Tunneling - The IPv6 packet is encapsulated inside an IPv4 packet.
 - Translation - Network Address Translation 64 (NAT64) allows IPv6-enabled devices to communicate with IPv4 devices.




IPv6 Addressing

- IPv6 Address Representation

- IPv6 Addresses:

- 128 bits in length
 - Every 4 bits is represented by a single hexadecimal digit
 - Hextet - unofficial term referring to a segment of 16 bits or four hexadecimal values.



IPv6 Addressing (Representation)

– Preferred format for IPv6 representation

2001	:	0DB8	:	0000	:	1111	:	0000	:	0000	:	0000	:	0200
2001	:	0DB8	:	0000	:	00A3	:	ABCD	:	0000	:	0000	:	1234
2001	:	0DB8	:	000A	:	0001	:	0000	:	0000	:	0000	:	0100
2001	:	0DB8	:	AAAA	:	0001	:	0000	:	0000	:	0000	:	0200
FE80	:	0000	:	0000	:	0000	:	0123	:	4567	:	89AB	:	CDEF
FE80	:	0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0001
FF02	:	0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0001
FF02	:	0000	:	0000	:	0000	:	0000	:	0001	:	FF00	:	0200
0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0001
0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0000	:	0000

IPv6 Addressing

- Rule 1 - Omit Leading 0s
 - In order to reduce or compress IPv6
 - First rule is to omit leading zeros in any hextet.

Preferred	2001:0DB8:0000:1111:0000:0000:0000:0200
No leading 0s	2001: DB8: 0:1111: 0: 0: 0: 200

Preferred	2001:0DB8:000A:1000:0000:0000:0000:0100
No leading 0s	2001: DB8: A:1000: 0: 0: 0: 100

Preferred	0000:0000:0000:0000:0000:0000:0000:0000
No leading 0s	0: 0: 0: 0: 0: 0: 0: 0

IPv6 Addressing


- Rule 2 - Omit All 0 Segments
 - A double colon (::) can replace any single, contiguous string of one or more 16-bit segments (hextets) consisting of all 0s.

Preferred	2001:0DB8:0000:0000:ABCD:0000:0000:0100
No leading 0s	2001: DB8: 0 : 0 : ABCD : 0 : 0 : 100
Compressed	2001:DB8::ABCD:0:0:100
or	
Compressed	2001:DB8:0:0:ABCD::100

Only one :: may be used.


Types of IPv6 Addresses

- IPv6 Address Types
 - Unicast- Single source IPv6 address.
 - Multicast - An IPv6 multicast address is used to send a single IPv6 packet to multiple destinations.
 - Anycast - An IPv6 anycast address is any IPv6 unicast address that can be assigned to multiple devices.




Types of IPv6 Addresses

- IPv6 Prefix Length
 - The IPv6 prefix length is used to indicate the network portion of an IPv6 address:
 - The prefix length can range from 0 to 128.
 - Typical IPv6 prefix length for most LANs is /64




Types of IPv6 Addresses

- IPv6 Unicast Addresses
 - Global Unicast - These are globally unique, Internet routable addresses.
 - Link-local - used to communicate with other devices on the same local link. Confined to a single link.
 - Unique Local - used for local addressing within a site or between a limited number of sites.



IPv6 Unicast Addresses

- IPv6 Link-Local Unicast Addresses
 - Enable a device to communicate with other IPv6-enabled devices on the same link only.
 - Are created even if the device has not been assigned a global unicast IPv6 address.
 - Are in the FE80::/10 range.
 - Note: Typically, it is the link-local address of the router that is used as the default gateway for other devices on the link.





IPv6 Unicast Addresses

- Structure of an IPv6 Global Unicast Address


- Currently, only global unicast addresses with the first three bits of 001 or 2000::/3 are being assigned

- A global unicast address has three parts:
 - Global routing prefix - network portion of the address that is assigned by the provider. Typically /48.
 - Subnet ID - Used to subnet within an organization.
 - Interface ID - equivalent to the host portion of an IPv4 address.



IPv6 Unicast Addresses


- Static Configuration of a Global Unicast Address
 - Router Configuration:
 - Similar commands to IPv4, replace IPv4 with IPv6
 - Command to configure and IPv6 global unicast on an interface is
ipv6 address ipv6-address/prefix-length



IPv6 Unicast Addresses (Static Conf.)

– Host Configuration:

- Manually configuring the IPv6 address on a host is similar to configuring an IPv4 address
- Default gateway address can be configured to match the link-local or global unicast address of the Gigabit Ethernet interface.



– Dynamic assignment of IPv6 addresses *:

- Stateless Address Autoconfiguration (SLAAC)
- Stateful DHCPv6

IPV6 ROUTING

- Configure IPv6 Static Routes
- Configure IPv6 Default Routes
- Configure EIGRP for IPv6

Configure IPv6 Static Routes


- The **ipv6 route** Command

```
Router(config)# ipv6 route ipv6-prefix/prefix-length {ipv6-address | exit-intf}
```

Parameter	Description
ipv6-prefix	Destination network address of the remote network to be added to the routing table.
prefix-length	Prefix length of the remote network to be added to the routing table.
ipv6-address	<ul style="list-style-type: none">Commonly referred to as the next-hop router's IP address.Typically used when connecting to a broadcast media (i.e., Ethernet).Commonly creates a recursive lookup.
exit-intf	<ul style="list-style-type: none">Use the outgoing interface to forward packets to the destination network.Also referred to as a directly attached static route.Typically used when connecting in a point-to-point configuration.


Configure IPv6 Static Routes

- **ipv6 unicast-routing** enables the router to forward IPv6 packets



Configure IPv6 Static Routes

- Next-Hop Options



R1#show ipv6 route

```
<output omitted>

C 2001:DB8:ACAD:1::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L 2001:DB8:ACAD:1::1/128 [0/0]
  via GigabitEthernet0/0, receive
C 2001:DB8:ACAD:4::/64 [0/0]
  via Serial0/0/0, directly connected
L 2001:DB8:ACAD:4::1/128 [0/0]
  via Serial0/0/0, receive
L FF00::/8 [0/0]
  via Null0, receive
R1#
```

R2#show ipv6 route

```
<output omitted>

C 2001:DB8:ACAD:2::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L 2001:DB8:ACAD:2::1/128 [0/0]
  via GigabitEthernet0/0, receive
C 2001:DB8:ACAD:4::/64 [0/0]
  via Serial10/0/0, directly connected
L 2001:DB8:ACAD:4::2/128 [0/0]
  via Serial10/0/0, receive
C 2001:DB8:ACAD:5::/64 [0/0]
  via Serial10/0/1, directly connected
L 2001:DB8:ACAD:5::2/128 [0/0]
  via Serial10/0/1, receive
L FF00::/8 [0/0]
  via Null0, receive
R2#
```


R3#show ipv6 route

```
<output omitted>

C 2001:DB8:ACAD:3::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L 2001:DB8:ACAD:3::1/128 [0/0]
  via GigabitEthernet0/0, receive
C 2001:DB8:ACAD:5::/64 [0/0]
  via Serial10/0/1, directly connected
L 2001:DB8:ACAD:5::1/128 [0/0]
  via Serial10/0/1, receive
L FF00::/8 [0/0]
  via Null0, receive
R3#
```

Configure IPv6 Static Routes

- Configure a Next Hop Static IPv6 Route
 - As with IPv4, must resolve the route to determine the exit interface to use to forward the packet
 - The IPv6 address matches the route for the directly connected network 2001:DB8:ACAD:4::/64 with the exit interface Serial 0/0/0.




```
R1#show ipv6 route
<output omitted>
C 2001:DB8:ACAD:1::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L 2001:DB8:ACAD:1::1/128 [0/0]
  via GigabitEthernet0/0, receive
C 2001:DB8:ACAD:4::/64 [0/0]
  via Serial0/0/0, directly connected
L 2001:DB8:ACAD:4::1/128 [0/0]
  via Serial0/0/0, receive
L FF00::/8 [0/0]
  via Null0, receive
R1#
```

```
R1# show ipv6 route
IPv6 Routing Table - default - 8 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route, B - BGP,
       R - RIP, H - NHRP, I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea,
       IS - ISIS summary, D - EIGRP, EX - EIGRP external, ND - ND Default,
       NDp - ND Prefix, DCE - Destination, NDr - Redirect, O - OSPF Intra,
       OI - OSPF Inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2, ON1 - OSPF NSSA ext 1,
       ON2 - OSPF NSSA ext 2

C 2001:DB8:ACAD:1::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L 2001:DB8:ACAD:1::1/128 [0/0]
  via GigabitEthernet0/0, receive
S 2001:DB8:ACAD:2::/64 [1/0]
  via 2001:DB8:ACAD:4::2
S 2001:DB8:ACAD:3::/64 [1/0]
  via 2001:DB8:ACAD:4::2
C 2001:DB8:ACAD:4::/64 [0/0]
  via Serial0/0/0, directly connected
L 2001:DB8:ACAD:4::1/128 [0/0]
  via Serial0/0/0, receive
S 2001:DB8:ACAD:5::/64 [1/0]
  via 2001:DB8:ACAD:4::2
L FF00::/8 [0/0]
  via Null0, receive
R1#
```

Configure IPv6 Static Routes


- Configure a Directly Connected Static IPv6 Route
 - Alternative to next hop is to specify the exit interface
 - Packet destined for 2001:DB8:ACAD:3::/64 network, forwarded out Serial 0/0/0 - no other lookups needed



```
R1# show ipv6 route
IPv6 Routing Table - default - 8 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route
      B - BGP, R - RIP, I1 - ISIS L1, I2 - ISIS L2
      IA - ISIS interarea, IS - ISIS summary, D - EIGRP, EX - EIGRP external
      ND - ND Default, NDp - ND Prefix, DCE - Destination, NDr - Redirect
      O - OSPF Intra, OI - OSPF Inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
C  2001:DB8:ACAD:1::/64 [0/0]
  via GigabitEthernet0/0, directly connected
L  2001:DB8:ACAD:1::1/128 [0/0]
  via GigabitEthernet0/0, receive
S  2001:DB8:ACAD:2::/64 [1/0]
  via Serial0/0/0, directly connected
S  2001:DB8:ACAD:3::/64 [1/0]
  via Serial0/0/0, directly connected
C  2001:DB8:ACAD:4::/64 [0/0]
  via Serial0/0/0, directly connected
L  2001:DB8:ACAD:4::1/128 [0/0]
  via Serial0/0/0, receive
S  2001:DB8:ACAD:5::/64 [1/0]
  via Serial0/0/0, directly connected
L  FF00::8 [0/0]
  via Null0, receive
R1#
```

Configure IPv6 Static Routes


- Configure a Fully Specified Static IPv6 Route
 - Fully specified static route must be used if IPv6 link-local address is used as next-hop



```
R1# show ipv6 route static | begin 2001:DB8:ACAD:2::/64
S 2001:DB8:ACAD:2::/64 [1/0]
  via FE80::2, Serial0/0/0
```

Configure IPv6 Static Routes

- Verify IPv6 Static Routes



```
R1# show ipv6 route 2001:db8:acad:3::/64
Routing entry for 2001:DB8:ACAD:3::/64
Known via "static", distance 1, metric 0
Route count is 1/1, share count 0
Routing paths:
  2001:DB8:ACAD:4::2
    Last updated 00:19:11 ago
R1#
```

```
R1# show running-config | section ipv6 route
ipv6 route 2001:DB8:ACAD:2::/64 2001:DB8:ACAD:4::2
ipv6 route 2001:DB8:ACAD:3::/64 2001:DB8:ACAD:4::2
ipv6 route 2001:DB8:ACAD:5::/64 2001:DB8:ACAD:4::2
R1#
```

Configure IPv6 Default Routes


- Default Static IPv6 Route

```
Router(config)# ipv6 route ::/0 (ipv6-address | exit-intf)
```

Parameter	Description
::/0	Matches any IPv6 prefix regardless of prefix length.
ipv6-address	<ul style="list-style-type: none">• Commonly referred to as the next-hop router's IPv6 address.• Typically used when connecting to a broadcast media (i.e., Ethernet).• Commonly creates a recursive lookup.
exit-intf	<ul style="list-style-type: none">• Use the outgoing interface to forward packets to the destination network.• Also referred to as a directly attached static route.• Typically used when connecting in a point-to-point configuration.


Configure IPv6 Default Routes

- Configure a Default Static IPv6 Route
 - R1 is a stub router because it is only connected to R2
 - More efficient to configure a default static IPv6 route in this topology



Configure IPv6 Default Routes


- Verify a Default Static Route
 - ::/0 mask indicates that none of the bits are required to match
 - If a more specific match does not exist, the default static IPv6 route matches all packets.



```
R1# ping 2001:0DB8:ACAD:3::1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2001:DB8:ACAD:3::1,
timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max
= 28/28/28 ms
R1#
```

Configure EIGRP for IPv6

```
R2# show running-config
<output omitted>
!
interface GigabitEthernet0/0
  ipv6 address 2001:DB8:CAFE:2::1/64
!
interface Serial0/0/0
  ipv6 address 2001:DB8:CAFE:A001::2/64
!
interface Serial0/0/1
  ipv6 address 2001:DB8:CAFE:A002::1/64
  clock rate 64000
!
interface Serial0/1/0
  ipv6 address 2001:DB8:FEED:1::1/64
```



```
R1# show running-config
<output omitted>
!
interface GigabitEthernet0/0
  ipv6 address 2001:DB8:CAFE:1::1/64
!
interface Serial0/0/0
  ipv6 address 2001:DB8:CAFE:A001::1/64
  clock rate 64000
!
interface Serial0/0/1
  ipv6 address 2001:DB8:CAFE:A003::1/64
```

```
R3# show running-config
<output omitted>
!
interface GigabitEthernet0/0
  ipv6 address 2001:DB8:CAFE:3::1/64
!
interface Serial0/0/0
  ipv6 address 2001:DB8:CAFE:A003::2/64
  clock rate 64000
!
interface Serial0/0/1
  ipv6 address 2001:DB8:CAFE:A002::2/64
```

Configure EIGRP for IPv6

```
R1(config)# interface s 0/0/0
R1(config-if)# ipv6 address fe80::1 ?
link-local Use link-local address

R1(config-if)# ipv6 address fe80::1 link-local
R1(config-if)# exit
R1(config)# interface s 0/0/1
R1(config-if)# ipv6 address fe80::1 link-local
R1(config-if)# exit
R1(config)# interface g 0/0
R1(config-if)# ipv6 address fe80::1 link-local
R1(config-if)#

```

```
R1# show ipv6 interface brief
GigabitEthernet0/0      [up/up]
FE80::1
 2001:DB8:CAFE:1::1
Serial0/0/0              [up/up]
FE80::1
 2001:DB8:CAFE:A001::1
Serial0/0/1              [up/up]
FE80::1
 2001:DB8:CAFE:A003::1
R1#


```

Same IPv6 link-local address is configured on all interfaces.

```
R2(config)# interface s 0/0/0
R2(config-if)# ipv6 address fe80::2 link-local
R2(config-if)# exit
R2(config)# interface s 0/0/1
R2(config-if)# ipv6 address fe80::2 link-local
R2(config-if)# exit
R2(config)# interface s 0/1/0
R2(config-if)# ipv6 address fe80::2 link-local
R2(config-if)# exit
R2(config)# interface g 0/0
R2(config-if)# ipv6 address fe80::2 link-local
R2(config-if)#

```


EIGRP for IPv6 Routing Domain



Configure EIGRP for IPv6

- The **ipv6 unicast-routing** global config mode command enables IPv6 routing on the router.
- Use the **ipv6 router eigrp** autonomous-system to enter EIGRP for IPv6 router configuration mode.
- Use the **eigrp router-id** **router-id** command is used to configure the router ID.
- By default, the EIGRP for IPv6 process is in a shutdown state and the **no shutdown** command is required to activate the EIGRP for IPv6 process.

```
R2(config)# ipv6 unicast-routing
R2(config)# ipv6 router eigrp 2
R2(config-rtr)# eigrp router-id 2.0.0.0
R2(config-rtr)# no shutdown
R2(config-rtr)#{
```




Configure EIGRP for IPv6

- Unlike EIGRP for IPv4 which uses the **network** command, EIGRP for IPv6 is configured directly on the interface using the **ipv6 eigrp autonomous-system** interface configuration command.

```
R2(config)# interface g 0/0
R2(config-if)# ipv6 eigrp 2
R2(config-if)# exit
R2(config)# interface s 0/0/0
R2(config-if)# ipv6 eigrp 2
R2(config-if)# exit
%DUAL-5-NBRCHANGE: EIGRP-IPv6 2: Neighbor FE80::1
(Serial0/0/0) is up: new adjacency
R2(config)# interface s 0/0/1
R2(config-if)# ipv6 eigrp 2
R2(config-if)#
```

```
R1(config)# interface g0/0
R1(config-if)# ipv6 eigrp 2
R1(config-if)# exit
R1(config)# interface s 0/0/0
R1(config-if)# ipv6 eigrp 2
R1(config-if)# exit
R1(config)# interface s 0/0/1
R1(config-if)# ipv6 eigrp 2
R1(config-if)#
```



The same passive-interface command used for IPv4 is used with EIGRP for IPv6.

```
R1(config)# ipv6 router eigrp 2
R1(config-rtr)# passive-interface gigabitethernet 0/0
R1(config-rtr)# end

R1# show ipv6 protocols

IPv6 Routing Protocol is "eigrp 2"
EIGRP-IPv6 Protocol for AS(2)
<output omitted>

Interfaces:
  Serial0/0/0
  Serial0/0/1
  GigabitEthernet0/0 (passive)
Redistribution:
  None
R1#
```

Configure EIGRP for IPv6

- Verifying EIGRP for IPv6
 - Use the **show ipv6 eigrp neighbors** command to view the neighbor table and verify that EIGRP for IPv6 has established an adjacency with its neighbors.
 - H - Lists the neighbors in order they were learned.
 - Address - IPv6 link-local address of the neighbor.
 - Interface - Local interface that received the Hello.
 - Hold - Current hold time.
 - Uptime - Time since this neighbor was added.
 - SRTT and RTO - Used by RTP.
 - Queue Count - Should always be zero.
 - Sequence Number - Used to track updates, queries, and reply packets.

R1# show ipv6 eigrp neighbors EIGRP-IPv6 Neighbors for AS(2)							
H	Address	Interface	Hold (sec)	Uptime (sec)	SRTT (ms)	RTO (ms)	Q Cnt Seq Num
1	Link-local address: FE80::3	Se0/0/1	13	00:37:17	45	270	0 8
0	Link-local address: FE80::2	Se0/0/0	14	00:53:16	32	2370	0 8

R1#

Neighbor's IPv6 Link-local Address.

Local Interface receiving EIGRP for IPv6 Hello packets.

Amount of time since this neighbor was added to the neighbor table.

Seconds remaining before declaring neighbor down.

The current hold time and is reset to the maximum hold time whenever a Hello packet is received.

Configure EIGRP for IPv6

- The **show ipv6 protocols** command displays the parameters and other information about the state of any active IPv6 routing protocol processes currently configured on the router.
 - EIGRP for IPv6 is an active dynamic routing protocol on R1.
 - These are the k values used to calculate the EIGRP composite metric.
 - The EIGRP for IPv6 router ID of R1 is 1.0.0.0.
 - Same as EIGRP for IPv4, EIGRP for IPv6 administrative distances have internal AD of 90 and external of 170 (default values).
 - The interfaces enabled for EIGRP for IPv6.

```
R1# show ipv6 protocols
IPv6 Routing Protocol is "connected"
IPv6 Routing Protocol is "ND"
IPv6 Routing Protocol is "eigrp 2"
EIGRP-IPv6 Protocol for AS(2) 1 Routing protocol and Process ID (AS Number)

Metric weight K1=1, K2=0, K3=1, K4=0, K5=0 2 K values used in composite metric

NSF-aware route hold timer is 240
Router-ID: 1.0.0.0 3 EIGRP Router ID
Topology : 0 (base)
  Active Timer: 3 min
  Distance: internal 90 external 170
  Maximum path: 16
  Maximum hopcount 100
  Maximum metric variance 1 4 EIGRP Administrative Distances

Interfaces:
  GigabitEthernet0/0
  Serial0/0/0
  Serial0/0/1 5 Interfaces enabled for EIGRP for IPv6
Redistribution:
  None
R1#
```

Configure EIGRP for IPv6

- Use the `show ipv6 route` command to examine the IPv6 routing table.
 - EIGRP for IPv6 routes are denoted with a D.
- The figure shows that R1 has installed three EIGRP routes to remote IPv6 networks in its IPv6 routing table:
 - 2001:DB8:CAFE:2::/64 via R3 (FE80::3) using its Serial 0/0/1 interface
 - 2001:DB8:CAFE:3::/64 via R3 (FE80::3) using its Serial 0/0/1 interface
 - 2001:DB8:CAFE:A002::/64 via R3 (FE80::3) using its Serial 0/0/1 interface

```
R1# show ipv6 route
<output omitted>

C 2001:DB8:CAFE:1::/64 [0/0]
    via GigabitEthernet0/0, directly connected
L 2001:DB8:CAFE:1::1/128 [0/0]
    via GigabitEthernet0/0, receive
D 2001:DB8:CAFE:2::/64 [90/3524096]
    via FE80::3, Serial0/0/1
D 2001:DB8:CAFE:3::/64 [90/2170112]
    via FE80::3, Serial0/0/1
C 2001:DB8:CAFE:A001::/64 [0/0]
    via Serial0/0/0, directly connected
L 2001:DB8:CAFE:A001::1/128 [0/0]
    via Serial0/0/0, receive
D 2001:DB8:CAFE:A002::/64 [90/3523840]
    via FE80::3, Serial0/0/1
C 2001:DB8:CAFE:A003::/64 [0/0]
    via Serial0/0/1, directly connected
L 2001:DB8:CAFE:A003::1/128 [0/0]
    via Serial0/0/1, receive
L FF00::/8 [0/0]
    via Null0, receive
R1#
```

Questions and Answers

